

Econ 2148, fall 2019  
Trees, forests, and causal trees

Maximilian Kasy

Department of Economics, Harvard University

## Agenda

- ▶ Regression trees: Splitting the covariate space.
- ▶ Random forests: Many trees.  
Using bootstrap aggregation to improve predictions.
- ▶ Causal trees: Predicting heterogeneous causal effects.  
Ground truth not directly observable, for cross-validation.

## Takeaways for this part of class

- ▶ Trees partition the covariate space and form predictions as local averages.
- ▶ Iterative splitting of partitions allows us to be more flexible in regions of the covariate space with more variation of outcomes.
- ▶ Bootstrap aggregation (bagging) is a way to get smoother predictions, and leads to random forests when applied to trees.
- ▶ Things get more complicated when we want to predict heterogeneous causal effects, rather than observable outcomes.
- ▶ This is because we do not directly observe a ground truth that can be used for tuning.



## Regression trees

- ▶ Suppose we have i.i.d. observations  $(X_i, Y_i)$  and want to estimate  $g(x) = E[Y|X = x]$ .
- ▶ Suppose we furthermore have a partition of regressor space into subset  $(R_1, \dots, R_M)$ .
- ▶ Then we can estimate  $g(\cdot)$  by averages in each element of the partition:

$$\hat{g}(x) = \sum_m c_m \cdot \mathbf{1}(x \in R_m)$$
$$c_m = \frac{\sum_i Y_i \cdot \mathbf{1}(X_i \in R_m)}{\sum_i \mathbf{1}(X_i \in R_m)}.$$

- ▶ This is a regression analog of a histogram.

## Recursive binary partitions



## Constructing the partition

- ▶ How to choose the partition?
- ▶ Start with the trivial partition with one element.
- ▶ Greedy algorithm (CART): Iteratively split an element of the partition, such that the in-sample prediction improves as much as possible.
- ▶ That is: Given  $(R_1, \dots, R_M)$ ,
  - ▶ For each  $R_m$ ,  $m = 1, \dots, M$ , and
  - ▶ for each  $X_j$ ,  $j = 1, \dots, k$ ,
  - ▶ find the  $x_{j,m}$  that minimizes the mean squared error, if we split  $R_m$  along variable  $X_j$  at  $x_{j,m}$ .
  - ▶ Then pick the  $(m, j)$  that minimizes the mean squared error, and construct a new partition with  $M + 1$  elements.
  - ▶ Iterate.

## Tuning and pruning

- ▶ Key tuning parameter: Total number of splits  $M$ .
- ▶ We can optimize this via cross-validation.
- ▶ CART can furthermore be improved using “**pruning.**”
- ▶ Idea:
  - ▶ Fit a flexible tree (with large  $M$ ) using CART.
  - ▶ Then iteratively remove (collapse) nodes.
  - ▶ To minimize the sum of squared errors,  
plus a penalty for the number of elements in the partition.
- ▶ This improves upon greedy search.  
It yields smaller trees for the same mean squared error.

## From trees to forests


- ▶ Trees are intuitive and do OK, but they are not amazing for prediction.
- ▶ We can improve performance a lot using either bootstrap aggregation (bagging) or boosting.
- ▶ **Bagging:**
  - ▶ Repeatedly draw bootstrap samples  $(X_i^b, Y_i^b)_{i=1}^n$  from the observed sample.
  - ▶ For each bootstrap sample, fit a regression tree  $\hat{g}^b(\cdot)$ .
  - ▶ Average across bootstrap samples to get the predictor

$$\hat{g}(x) = \frac{1}{B} \sum_{b=1}^B \hat{g}^b(x).$$


- ▶ This is a technique for smoothing predictions.  
The resulting predictor is called a “random forest.”
- ▶ Possible modification:  
Restrict candidate splits to random subset of predictors in each tree-fitting step.




## An empirical example (courtesy of Jann Spiess)




# OLS



## Regression tree



# Random forest



## Causal trees

- Suppose we observe i.i.d. draws of  $(Y_i, D_i, X_i)$ , and wish to estimate

$$\tau(x) = E[Y|D = 1, X = x] - E[Y|D = 0, X = x].$$

- Motivation: This is the conditional average treatment effect under an unconfoundedness assumption on potential outcomes,

$$(Y^0, Y^1) \perp D | X.$$

- This is relevant, in particular, for targeted treatment assignment.
- We might, for a given partition  $\mathcal{R} = (R_1, \dots, R_M)$ , use the estimator

$$\hat{\tau}(x) = \sum_m (c_m^1 - c_m^0) \cdot \mathbf{1}(x \in R_m)$$
$$c_m^d = \frac{\sum_i Y_i \cdot \mathbf{1}(X_i \in R_m, D_i = d)}{\sum_i \mathbf{1}(X_i \in R_m, D_i = d)}.$$

## Targets for splitting and cross-validation

- ▶ Recall that CART uses greedy splitting.  
It aims to minimize in-sample mean squared error.
- ▶ For tuning, we proposed to use the out-of-sample mean squared error in order to choose the tree depth.
- ▶ Analog for estimation of  $\tau(\cdot)$ : Sum of squared errors,

$$SSE(\mathcal{S}) = \sum_{i \in \mathcal{S}} ((\tau_i - \hat{\tau}(X_i))^2 - \tau_i^2),$$

where  $\mathcal{S}$  is either the estimation sample, or a hold-out sample for cross-validation.  
(The term  $\tau_i^2$  is added as a convenient normalization.)

- ▶ Problem:  $\tau_i$  is not observed.

## Targets continued

- Solution: We can rewrite  $SSE(\mathcal{S})$ ,

$$SSE(\mathcal{S}) = \sum_{i \in \mathcal{S}} (\hat{\tau}(X_i, \mathcal{R}) \cdot (\hat{\tau}(X_i, \mathcal{R})^2 - 2\tau_i)).$$

- Suppose we split our sample into  $(\mathcal{S}^1, \mathcal{S}^2)$ , use  $\mathcal{S}^1$  for estimation, and  $\mathcal{S}^2$  for tuning. Let  $\hat{\tau}_j(X, \mathcal{R})$  be the estimator based on sample  $\mathcal{S}^j$ .
- An unbiased estimator of  $SSE(\mathcal{S}^2)$  (for tuning) is then given by

$$\widehat{SSE}(\mathcal{S}^2) = \sum_{i \in \mathcal{S}} (\hat{\tau}_1(X_i, \mathcal{R}) \cdot (\hat{\tau}_1(X_i, \mathcal{R})^2 - 2\tau_2(X_i, \mathcal{R}))).$$

- An analog to the in-sample sum of squared errors (for CART splitting) is given by

$$\widehat{SSE}(\mathcal{S}^1) = \sum_{i \in \mathcal{S}} (-\hat{\tau}_1(X_i, \mathcal{R})^2).$$

## References

- ▶ *Friedman, J., Hastie, T., and Tibshirani, R. (2001). The elements of statistical learning, volume 1. Springer series in statistics Springer, Berlin, chapters 8 and 9.*
- ▶ *Athey, S. and Imbens, G. (2016). Recursive partitioning for heterogeneous causal effects. Proceedings of the National Academy of Sciences, 113(27):7353–7360.*